

Stablecoin StressBench: An Execution-Aware Benchmark for Stablecoin Dislocations

Anonymous Author(s)

Abstract

Quoted stablecoin price gaps are not trades: we call dislocations where a quoted basis exceeds a threshold *optical*, reserving *executable* for the subset where net profit survives VWAP slippage, taker fees, and settlement latency after routing through a live order book. Building **Stablecoin StressBench** on the March 2023 USDC/SVB event, we find that 34.3% of minutes show a > 10 bps optical dislocation but only 2.88% are executable: a **12 $\times$ price-to-execution gap** (block bootstrap 95% CI: $8\times$ – $24\times$), with no calm-trained model closing it and the hindsight oracle exceeding +225 bps net on the executable arbitrage task. Cross-mechanism meta-labeling trained on Terra/LUNA achieves +82.5 bps on the SVB test split (50.8% of oracle returns), demonstrating that order-book execution microstructure transfers across structurally distinct stress mechanisms. A three-way RL diagnostic isolates why: unconditioned PPO-GRU degenerates to NoTrade at 2.3% positive density; conditioned PPO-GRU at 17% density trades 919 times but earns -29.2 bps; only meta-labeling’s explicit positive-label supervision reaches +82.5 bps—establishing that density is necessary but not sufficient.

Keywords

stablecoin arbitrage, market microstructure, execution-aware labels, financial benchmarks, cryptocurrency

ACM Reference Format:

Anonymous Author(s). 2026. Stablecoin StressBench: An Execution-Aware Benchmark for Stablecoin Dislocations. In *Proceedings of ACM International Conference on AI in Finance (ICAIF ’26)*. ACM, New York, NY, USA, 7 pages.

1 Introduction

During stablecoin stress events, peg fragmentation creates cross-venue price gaps that appear to be arbitrage opportunities. In practice, VWAP slippage, taker fees, and CEX-to-CEX settlement latency jointly eliminate most apparent profit. We call dislocations where the quoted basis exceeds a threshold *optical*, and those where net profit survives execution costs *executable*.

Prior machine-learning work on stablecoin predictability uses price-based labels to define “profitable” windows, which conflates the optical and executable layers and systematically overstates how predictable actual trading opportunities are. **Stablecoin StressBench** is designed to close this gap. It provides per-minute execution-aware labels that separate three dislocation layers: (1) an *optical* dislocation (quoted basis exceeds a threshold), (2) an *executable* dislocation (net profit is positive after VWAP execution, fees, and settlement at a specified notional), and (3) a *predictable* dislocation (a model correctly identifies the window *ex ante*). The distance between these layers is the paper’s core measurement.

We build the benchmark on the March 2023 USDC/SVB stress event, where Silicon Valley Bank’s collapse triggered a brief USDC depeg. Using Binance L2 order-book data¹ and a full VWAP book-walking model, we find that 34.3% of 1-minute windows show a > 10 bps price basis but only 2.88% remain net-profitable at \$10K notional: a 12 $\times$ price-to-execution gap. The headline result is robust to buy-leg proxy error: a 20% depth haircut moves the executable rate from 2.88% to 2.40% and widens the ratio from 12 $\times$ to 14 $\times$, with the oracle remaining above +130 bps. We then construct a model ladder spanning rule-based thresholds, statistical baselines, gradient-boosted classifiers, sequence models (GRU), and meta-labeling filters, evaluating each model’s ability to identify the 2.88% of executable windows *ex ante*. Every model trained on calm-control data loses money on executable arbitrage tasks; the hindsight oracle gap exceeds 260 bps. Cross-mechanism meta-labeling, trained on a different stress event (Terra/LUNA 2022), is the only approach that achieves positive net returns (+82.5 bps, 50.8% oracle capture).

Contributions.

- (1) **Execution-aware benchmark and labels:** per-minute labels with documented order-book depth provenance, VWAP book-walking, fees, and a hindsight oracle ceiling; accompanied by an 18-event stress taxonomy (seven mechanism classes, 2020–2023) that functions as data-tier governance specifying permitted claims per event, and an oracle gap decomposition into FP cost, timing, and sizing components. (§4, §3.3; Tables 4, 3, 5)
- (2) **Cross-mechanism transfer:** meta-labeling trained on Terra/LUNA achieves +82.5 bps (50.8% oracle capture) on the SVB test split—to our knowledge the first demonstration that order-book execution microstructure transfers across structurally distinct stress mechanisms, on the single cross-mechanism pair available with route-complete test data. (§6.5, Table 8)
- (3) **RL diagnostic—density necessary, supervision sufficient:** conditioned PPO-GRU at 17% positive density earns -29.2 bps vs. meta-labeling’s +82.5 bps at identical density, isolating explicit positive-label supervision as the additional requirement beyond RL reward optimisation. (§6.8, Table 10)

2 Related Work

The literature most relevant to our work spans stablecoin stress events, limits to arbitrage, and financial machine learning.

Prior work on stablecoin stress provides the empirical backdrop for our benchmark. Uhlig [1] analyses the Terra/LUNA collapse, whose algorithmic redemption spiral provides a structurally distinct trigger from the SVB bank shock, making it a clean out-of-distribution test for cross-mechanism transfer. Griffin and Shams [2] document that quoted prices can reflect non-fundamental demand,

ICAIF ’26, Milan, Italy
2026.

¹The BTC-USDC buy leg uses kline-proxy depth for the SVB window; all other legs use real L2. The 20% depth haircut sensitivity in §6.1 bounds this approximation.

motivating our separation of optical from executable dislocations. Hernandez Cruz et al. [21] study USDC-USDT DEX liquidity during the SVB collapse and confirm the importance of execution-layer analysis.

The limits-to-arbitrage literature establishes the theoretical foundations that our labels operationalise. Shleifer and Vishny [4] and Gromb and Vayanos [5] show that capital constraints and execution frictions bound arbitrage even when price gaps are observed. Makarov and Schoar [7] demonstrate that cross-venue crypto price gaps are not executable profits once latency and book depth are accounted for, directly motivating our 12 $\times$ price-to-execution gap. Hautsch et al. [3] quantify settlement latency as a first-order cost, which we implement as the fixed $\delta=5$ bps penalty; Cont and Kukanov [6] establish VWAP execution price as the correct profitability object, which our book-walking model in §4.2 implements directly. Gogol et al. [18] introduce a Maximal Arbitrage Value between a USDC-ETH AMM pool and Binance; our oracle extends this to route-level CEX simulation and model-evaluation benchmarking.

The financial machine learning literature informs both our methodology and our positioning. López de Prado [8] introduces meta-labeling as a technique for filtering a primary signal with a secondary model that learns the conditions under which the primary signal is reliable. We extend this architecture to the cross-mechanism setting in §6.5, training on Terra/LUNA and evaluating on SVB—a transfer that, to our knowledge, has not been attempted before. Cintra and Holloway [9] apply BOCPD [15] to detect the Terra depeg from Curve AMM pool reserves, and demonstrate cross-event transfer via AMM depletion signals; however, their signals differ fundamentally from CEX cross-quote basis, which is noisier and mean-reverting, and no oracle evaluation is reported. On reinforcement learning, Moallemi and Wang [16] and Ning et al. [17] establish the finite-horizon MDP framing for optimal execution that we adopt in §5.2. Mohl et al. [22] demonstrate the feasibility of GPU-scale multi-agent RL on production limit-order-book data, training IPPO on approximately 400 million orders for a joint execution and market-making task. Olby et al. [23] construct a reactive Simudyne simulator that decomposes execution costs into market impact and execution risk, benchmarking RL timing policies against an Almgren–Chriss efficient frontier. Both works confirm that RL can learn effective execution policies when the environment is reactive; our benchmark contributes the complementary finding that in non-reactive historical replay with structurally rare positive labels, positive-label density and explicit supervision are necessary preconditions. No prior stablecoin-stress benchmark provides one-minute executable labels, route-level depth and fee modelling, a hindsight oracle ceiling, and cross-mechanism transfer evaluation jointly; StressBench is the first to do so.

3 Benchmark Design

3.1 Data Sources and Instruments

The final dataset contains 56,134 1-minute bars across six windows (Table 2). Price features come from Binance Vision (klines, aggTrades) and Coinbase REST candles. Net-profit labels use Binance BTC-USDC and BTC-USDT depth plus Coinbase BTC-USD as the reference leg.

Table 1: Route-leg depth provenance.

Leg	SVB Mar 2023	2024 windows	
Buy BTC-USDC	kline-proxy [†]	real L2 (futures)	† BTC-USDC
Cross USDC-USDT	real L2 (Mar 12–20)	real L2 (futures)	
Sell BTC-USDT	real L2 (futures)	real L2 (futures)	
Ref BTC-USD	candle-proxy	candle-proxy	

perpetual listed Jan 2024; all 2022 windows kline-proxy.

Table 2: Dataset splits (56,134 1-min bars).

Split	Event	Min.	Pos.
Train	Calm control ($\times 3$)	28,776	0.00%
Validation	Terra/LUNA May 2022	11,526	2.30%
Test	USDC/SVB Mar 2023	15,832	2.88%




Figure 1: Cross-quote basis (bps) during the SVB test window (Mar 10–20 2023). USDC basis (navy): 12.4% of minutes > 10 bps; USDT basis (red): 27.4%; max-abs (gold): 34.3%. The price-to-execution gap is 12 $\times$.

Every depth record is tagged with a provenance field (Table 1). The sell leg (BTC-USDT) uses real L2 futures bookDepth for all 2023–2024 windows. The cross leg (USDC-USDT) uses real L2 from 2023-03-12 onward. The buy leg (BTC-USDC) uses kline-proxy for the SVB window (BTC-USDC perpetual did not exist on Binance until 2024) and real L2 for 2024 windows. To bound this proxy error, we apply a 20% depth haircut to the kline-proxy buy leg: the executable rate falls from 2.88% to 2.40%, the ratio widens from 12 $\times$ to 14 $\times$, and the oracle stays above +130 bps. The headline result is robust to buy-leg proxy error.

3.2 Event-Based Split Design

Splits are defined by event identity (Figure 1). No model trained on the train split has seen stress dynamics; threshold calibration on validation uses an independent stress event. Labels are constructed by aligning each bar’s outcome to the observation window ending at $t+h$, with explicit checks confirming no future information leaks into the input features.

3.3 Historical Stress Taxonomy: Data-Tier Governance

The 18-event taxonomy (2020–2023) is a benchmark governance artefact: it specifies which events carry sufficient data quality for execution-grade claims, and which are restricted to price magnitude estimates or mechanism classification only. It is shipped with the benchmark to foreclose over-interpretation of results, not as an empirical finding in its own right. Seven mechanism classes are covered: *algorithmic/reflexive* (FEL, IRON/TITAN, MIM, UST, USDD); *fiat-reserve bank shock* (USDC/SVB); *regulatory winddown* (BUSD);

Table 3: Data tiers, mechanisms, and execution frictions.

Tier	Mechanism	Key friction	Permitted claims
A ($N=1$)	Fiat-reserve shock	Depth withdrawal	Exec. gap; oracle gap; P&L
B ($N=11$)	Alg., exchange, DeFi	AMM venue seg.	slippage; Price magnitude (est.)
C ($N=5$)	RWA, niche	Thin market	Taxonomy only

exchange credit (Celsius/3AC, HUSD, FTX [10]); *DeFi pool imbalance* (Curve/UST, USDT/Curve, USDC/DAI); *collateral liquidation* (DAI Black Thursday); *RWA/niche* (Acala aUSD, USDR).

Data availability governs permitted claims (Table 3). Tier-A events (USDC/SVB test) support execution-gap, oracle-gap, and model evaluation claims. Tier-B ($N=11$, OHLCV/pool data) support price magnitude estimates only. Tier-C ($N=5$, partial/DEX-only) support mechanism taxonomy only.

3.4 Benchmark Versioning

StressBench v1.0 ships with the SVB test split; near-term Tier-A expansion targets are the USDC recovery window (Mar 15–31 2023) and 2024 BTC-USDC windows where real-L2 eliminates the kline-proxy. Release format: versioned Hugging Face dataset under CC-BY 4.0.

3.5 Feature Sets

Four nested feature sets isolate the marginal information value of

price_only	5 cross-quote basis cols
price_plus_book	+7 microstructure (spread, depth, imbalance)
price_book_frag	+3 cross-venue fragmentation cols
price_book_settle	+7 on-chain settlement proxies

4 Execution-Aware Labels

Labels operationalise the three-layer framework introduced in §1. The price-to-execution gap measures Layer 1 divided by Layer 2. The oracle gap measures Layer 3 against Layer 2 in hindsight. This section details label construction.

4.1 Cross-Quote Basis and Labels

$$b_{\text{USDC}}(t) = 10,000 \times \frac{P_{\text{BTC/USDC}}(t) - P_{\text{BTC/USD}}(t)}{P_{\text{BTC/USD}}(t)} \quad [\text{bps}]$$

$b_{\text{max}} = \max(|b_{\text{USDC}}|, |b_{\text{USDT}}|)$. Basis label: $y_{\text{basis}}(t) = 1[|b_{\text{USDC}}(t+h)| > \tau]$. Primary task: basis_usdc_1m_gt10bps.

4.2 VWAP Execution and Net Profit Labels

$$\text{net}(q, t) = 10,000 \times \left[\frac{\text{VWAP}_{\text{sell}} - \text{VWAP}_{\text{buy}}}{P_{\text{ref}}} - f_{\text{buy}} - f_{\text{sell}} - \delta_{\text{settle}} \right] \quad [\text{bps}] \quad (1)$$

$$y_{\text{arb}}(q, h, t) = 1 \left[\max_{s \in [t, t+h]} \text{net}(q, s) > c \right] \quad (2)$$

Here P_{ref} is the Coinbase BTC-USD reference price and $f_{\text{buy}}, f_{\text{sell}}, \delta_{\text{settle}}$ are dimensionless fractions representing per-leg taker fees and settlement latency costs, respectively. We set the profit threshold $c=10$ bps, taker fee to 4 bps per side, and $\delta=5$ bps; the primary economic task is executable_arb_q10000_5m. The hindsight oracle trades every minute where $\text{net}(q, t) > c$, establishing a

performance ceiling that no deployable model—which must decide without future information—can exceed.

5 Models and Evaluation

5.1 Model Ladder

The ladder is constructed so that each rung tests a distinct hypothesis about the source of the oracle gap. Economic anchors—NoTrade at 0 bps and the hindsight Oracle—bracket the achievable range. Price-threshold rules (PriceBasis{10, 25}bps, GrossArbThreshold) test whether simple signal filters suffice. Tabular classifiers (logistic regression, Lasso, Random Forest, XGBoost, LightGBM) trained on four nested feature sets test whether microstructure conditioning closes the gap. An expected-net-profit regressor that directly targets $\hat{y}=\text{net}(q, t)$ tests whether the gap is a label-mismatch artefact; failure here implies it is structural. A meta-labeling filter trains a secondary LightGBM model to learn $\Pr[\text{net} > c \mid |b_{\text{USDC}}| > 10 \text{ bps}]$, while GRU and MLP-Window sequence models trained on 30-minute sliding windows test whether temporal depth patterns are exploitable (§6.4).

5.2 RL Execution Policy Baseline

We frame the entry-timing problem as a finite-horizon MDP. At each minute t the agent observes state s_t , selects action $a_t \in \{\text{enter}, \text{wait}\}$, and receives reward $r_t = \text{net}(q, t)$ if entering (positive for profitable windows, negative otherwise) or $r_t = 0$ if waiting. The state s_t is the 30-minute look-back window of price_plus_book features, the smallest feature set that delivers meaningful microstructure lift over price-only signals (see §6.5 for SHAP evidence). This feature set is available in real time without settlement-latency dependency. We train a PPO agent with a GRU policy network (one hidden layer, 64 units, 30-step look-back) using the cross-mechanism protocol: training on Terra/LUNA stress windows only, evaluation on the USDC/SVB test split without retraining. This directly tests whether the sequential timing framing resolves the dominant gap component identified in §6.7.

5.3 Calibration and Evaluation

Threshold calibration on validation split: $\theta^* = \arg \max_{\theta} \sum_{i: \hat{p}_i > \theta} \text{net}(q, t_i)$ subject to ≥ 25 trades. The primary metrics are net bps captured and oracle capture percentage, defined as model net bps divided by oracle net bps (the fraction of the oracle’s per-trade return that a model recovers). Secondary metrics are AUROC and AUPRC. NoTrade at 0 bps is the economic floor.

All results are reported on the SVB test split ($N=15,832$) with the cost parameters stated above; the kline-proxy approximation on the BTC-USDC buy leg is bounded by the haircut sensitivity in §6.1.



6 Results

6.1 Price Gaps Are Common, Execution Is Rare

At 10 bps, 34.3% of test minutes exceed the max-basis threshold; only 2.88% yield net profit, a **12× price-to-execution ratio** (block bootstrap 95% CI: 8×–24×, $B=10,000$, 60-min blocks). The gap persists at all tested thresholds (Table 4) and is robust across cost

Table 4: Price-to-execution gap (\$10K, 1-min horizon, test split).

Thresh.	$P_{\max}$	P_{USDC}	Exec.	Ratio
0 bps	95.86%	82.96%	3.34%	29×
5 bps	61.99%	25.32%	3.03%	20×
10 bps	34.33%	12.45%	2.88%	12×
25 bps	9.54%	6.48%	2.62%	3.6×

**Figure 2: Oracle gap decomposition (\$10K notional). Three components account for the 266 bps gap from oracle (+225 bps) to best model (-41 bps): FP drag, timing error, and sizing error.****Figure 3: Notional scaling (SVB test split). The executable rate collapses from 2.88% at \$10K to 0.03% at \$500K: book depth, not fees, is the binding institutional constraint.**


regimes: at the worst-case setting (10 bps settlement, high fees), the executable rate falls to 2.40% and the ratio widens to 14×. At \$50K notional the executable rate falls to 1.62%; at \$500K it reaches 0.03%, confirming that book depth, not fees, prices out institutional participants. At institutional scale, book depth rather than fees is the binding constraint (0.03% executable rate at \$500K), consistent with the depth-based limits-to-arbitrage arguments in §2.

Claim boundaries and data-tier scope are stated in §5.3; the Terra/LUNA cross-mechanism check (5.9× ratio; §6.6) provides directional evidence that the gap is not SVB-specific.

Table 5: Oracle gap summary (USDC/SVB test split[§]). * =only non-oracle tabular result above zero; ‡ =sequence model (30-min window).

Task / Model	Oracle	Best model	Gap	AUPRC
basis_usdc_1m	+162.2 bps	+17.2 bps*	145 bps	0.253
exec_arb_q10k_5m	+225.0 bps	-41.1 bps	266 bps	0.060
exec_arb_q50k_5m	+147.1 bps	-113.0 bps	260 bps	0.063
GRU (seq) [‡] , basis_usdc_1m	+162.2 bps	-239.0 bps	401 bps	—

*lgbm@price_book_settle, 106 trades. ‡GRU 30-min window (§6.4). §Kline-proxy buy leg; 20% haircut sensitivity in §6.1.

**Figure 4: Running mean net bps per trade (USDC/SVB test split). Oracle (gold) +162.2 bps; cross-mech meta-label (green) +82.5 bps; LightGBM (navy) +17.2 bps; PriceBasis10bps (red) -269.5 bps; NoTrade (grey dashed) 0 bps. Shaded region: SVB shock window.**

6.2 Oracle Gap: Profitable Windows Exist but Only in Hindsight

The oracle earns +162.2 bps on the basis task and +225.0 bps on the executable arbitrage task (Table 5); Figure 4 shows the full performance hierarchy.

6.3 The Model Ladder: Structural Oracle Gap

Table 6 shows the full model ladder with loss attribution. Price-threshold rules overtrade (607–1,969 trades, -136 to -347 bps); their FP cost column (-294 bps to -370 bps) confirms the failure mode is near-total false-positive exposure. LightGBM achieves +17.2 bps (106 trades, AUROC 0.72), the only non-oracle result above zero, but 145 bps below the oracle. The ExpNetProfitRegressor reaches -73.8 bps despite directly targeting the economic criterion: the gap is structural, not a label-mismatch artefact.

Ranking accuracy does not imply positive returns. At 2.88% positives, precision-recall curves expose the imbalance problem directly (Figure 5). PriceBasis10bps achieves AP 0.55 and AUROC 0.83 while losing -270 bps; LightGBM achieves AP 0.25 and AUROC 0.72 as the sole non-oracle model above zero. High AP or AUROC is necessary but not sufficient for positive returns. High AUROC with low Hit% reflects the threshold regime: the calibrated threshold sits in a high-recall, low-precision region where most trades are false positives, even though the model's overall ranking is correct.

Table 6: Tabular model ladder (basis_usdc_1m_gt10bps, test split). Hit%: fraction of model trades that were net-profitable. [†]ceiling; *only result above zero. Sequence and RL models in Tables 7 and 10.

Model	Feat.	Net bps	Trades	AUROC	Mean bpswhen wrong	Hit%
Oracle [†]	—	+162.2	315	—	—	100%
lgbm*	all	+17.2	106	0.717	−65.2	33%
logistic	all	−75.8	648	0.143	−75.8	0%
gross_arb	px	−136.4	607	0.531	−150.9	4%
price_10	px	−270.0	1969	0.833	−294.6	6%
price_25	px	−347.3	1025	0.746	−370.7	5%
no_trade	—	0.0	0	—	—	—
ExpNetProfit	px	−73.8	537	—	−89.1	3%

px = price_only; all = price_book_settle.




Figure 5: Precision-recall curves (basis_usdc_1m_gt10bps, SVB test split). High AP does not imply positive net bps: PriceBasis achieves AP 0.55 while losing −270 bps.

6.4 Sequence Models: Confirming the Structural Oracle Gap

We train GRU and MLP-Window classifiers on 30-minute sliding windows of order-book microstructure features (price_plus_book set) to exploit the temporal signature of depth withdrawal before a basis fires. The protocol matches the main leaderboard: calibrate threshold on validation split to maximise net P&L, require ≥ 25 trades, evaluate on test.

Table 7 reports results on basis_usdc_1m_gt10bps.

Table 7: Sequence model results (basis_usdc_1m_gt10bps, SVB test data, test split). Oracle gap = +162.2 bps oracle minus model net bps.

Model	Net bps	Trades	AUROC	Gap
GRU (30-min)	−239.0	656	0.798	401 bps
MLP-Window (30-min)	−340.7	1,052	0.648	502 bps

Calm-control trained (0% positive examples); threshold calibrated on validation.

Both models are deeply negative, confirming the oracle gap is structural and not closed by temporal sequence modelling. AUROC 0.80 for GRU combined with −239bps net illustrates the

ranking-vs-returns disconnect: the calibrated threshold admits twice as many trades as the oracle because calm training produces zero positive stress examples. Cross-mechanism meta-labeling (§6.5) bypasses this by training on Terra/LUNA stress dynamics, achieving +82.5 bps.

6.5 Cross-Mechanism Transfer: Benchmark Generalisation Evidence

When the meta-labeling filter is trained on the calm-control split, it yields zero meta-positives by construction (0% positive rate). We therefore retrain on the Terra/LUNA validation split, which contains 1,559 primary fires and 265 meta-positives (17.0% positive rate), and evaluate cross-mechanism on the SVB test split. All fitting and threshold selection use Terra/LUNA only; no SVB labels are observed before final evaluation. With price_plus_book features the filter achieves +82.5 bps (397 trades, 50.8% oracle capture) vs. zero with calm training (Table 8). This is consistent with execution microstructure that transfers across stress mechanisms: the same order-book features that predict profitable windows during an algorithmic collapse (Terra/LUNA) also predict them during a fiat-reserve shock (SVB). The meta-model’s dominant features mirror the main SHAP ranking: ask-side depth at 10 bps and bid-ask spread are the leading microstructure signals in both calm-trained LightGBM and the Terra-trained meta-model secondary classifier. The structural explanation: CEX market makers withdraw depth under uncertainty regardless of trigger, because the cause is opaque at the order-book timescale, making depth withdrawal a consistent cross-mechanism precursor whose feature signature transfers while the magnitude differs by event (5.9× Terra/LUNA vs. 12× SVB). We are not aware of prior work demonstrating cross-mechanism transfer for execution-aware stablecoin models.

Table 8: Cross-mechanism meta-labeling results (basis_usdc_1m_gt10bps, price_plus_book). Oracle capture = meta net bps ÷ oracle net bps.

Training split	Net bps	Trades	Oracle cap.
Oracle (ceiling)	+162.2	315	100.0%
Terra/LUNA (cross-mech.)	+82.5	397	50.8%
Calm-control (baseline)	0.0	0	0.0%

signal $|b_{USDC}| > 10$ bps; threshold calibrated on validation.

Regime detectors, false-positive diagnosis, and SHAP. Three regime-detection filters were evaluated as two-stage pre-filters. CUSUM achieves near-perfect recall (0.9995) but at an 83.5% false alarm rate; EWMA reaches 86.9% accuracy with 0.95% recall and a 0.65% false alarm rate. BOCPD [15] achieves AUROC 0.229 on CEX cross-quote data, which contrasts with Cintra and Holloway [9], who detected the Terra depeg 12 hours early from Curve pool reserves: AMM reserves accumulate gradually, whereas CEX cross-quote basis is mean-reverting, making changepoint detection unreliable on exchange prices.

A structural diagnosis of false positives reveals that the dominant failure mode of price-threshold rules is route-direction mismatch rather than thin books. PriceBasis10bps produces 2,002 trades, of which 581 (29%) are false positives at a mean cost of −38.7 bps each.

Table 9: Cross-mechanism comparison.

Metric	USDC/SVB (Tier-A)	Terra/LUNA (val.)
Minutes	15,832	11,526
Mechanism	Fiat-reserve	Alg./reflexive
Price > 10 bps	34.3%	13.5% [†] Kline-proxy
Executable	2.88%	2.30%
P-to-E ratio	12×	5.9×
Oracle bps	+225.0 bps	+88 bps [†]

labels; directional only.

True-positive windows carry a mean USDC basis of 344 bps, reflecting a genuine discount, whereas false-positive windows carry just 0.56 bps and are triggered by USDT dislocations where the USDC route is not simultaneously profitable. Notably, false-positive windows exhibit *higher* average bid depth (72,805 vs. 59,961 BTC units) than true positives, ruling out liquidity as the cause. LightGBM’s 33% hit rate (compared to 6% for the rule) demonstrates that microstructure conditioning partially addresses this mismatch, though a 67% false-positive rate persists.

SHAP analysis identifies price features as carrying 94.8% of total importance in LightGBM, with ask-side depth and bid-ask spread contributing just 3.4% combined. The causal interpretation is straightforward: a model trained on calm-market data learns calm price dynamics, not the depth-withdrawal signature of stress events, so the microstructure features—while present—are rendered uninformative by the training distribution. This explains why the oracle gap is structural rather than a label-mismatch artefact, and why cross-mechanism training on Terra/LUNA resolves it: the meta-model encounters stress microstructure at 17% positive density and learns the depth-withdrawal signature that calm training cannot recover.

6.6 Terra/LUNA Replication: Pattern Generality

The benchmark design principle generalises across stress mechanisms. The Terra/LUNA validation split ($N=11,526$, May 2022) uses kline-proxy depth (Tier-B) but replicates the core pattern: 13.5% price-basis > 10 bps, only 2.30% executable, a 5.9×

AMM or exchange-credit events would require route-complete L2 archives for those venues, which are not currently available.

On the model side, the conditioned RL result (−29.2 bps, 919 trades) may reflect overfitting to Terra/LUNA’s specific depth-withdrawal timing rather than a fundamental limitation of sequential optimisation; disentangling these alternatives requires additional cross-mechanism pairs or extended training. A further constraint is the absence of market-impact feedback: unlike reactive-simulator approaches such as Olby et al. [23], where agents observe how their own order flow depletes book depth, the PPO-GRU is trained on historical replay and cannot model this loop. In a thin-book stress regime this systematically miscalibrates entry sizing, and we regard it as a plausible contributor to the 919-trade overtrading pattern at identical training density to meta-labeling. Integrating the cross-mechanism training protocol developed here into a reactive simulator represents the most direct path to resolving the timing component of the oracle gap.

8 Conclusion

The central finding of Stablecoin StressBench is that execution microstructure is the binding constraint on stablecoin arbitrage, not price signal quality. A 34.3% optical dislocation rate collapses to 2.88% once order-book depth and transaction costs are applied. No model trained on calm-market data closes this gap. Cross-mechanism training on a different stress event is the empirically validated path to positive net returns (+82.5 bps, 50.8% oracle capture). No calm-control-trained model earns positive returns on the executable arbitrage task; LightGBM (+17.2 bps, basis task only) is the sole non-oracle positive result, 145 bps below the oracle ceiling. The oracle gap decomposition (§6.7) identifies false-positive cost as the dominant quantifiable drag and timing as the primary structural barrier.

The RL experiments (§6.8) sharpen the benchmark implication: **positive example density is necessary but not sufficient; explicit supervision of the positive subset is the additional requirement.** The unconditioned PPO-GRU degenerates (2.3% density too sparse). The conditioned PPO-GRU trained on the same 17% primary-signal windows as the meta-label does trade (−29.2 bps, 919 trades) but cannot match meta-labeling (+82.5 bps), leaving an 111.7 bps gap attributable to meta-labeling’s explicit positive-label supervision rather than RL reward optimisation.

To support reproducibility, the anonymised repository (URL redacted for review) provides per-minute executable labels for all six benchmark windows, the VWAP book-walking label-construction code with route-level depth provenance, the four nested feature sets, the hindsight oracle implementation and model-ladder harness, and the cross-mechanism meta-labeling and PPO-GRU training scripts. Calm-control and SVB windows are reproducible from the public Binance Vision archive; route-complete L2 requires a Tardis subscription.

Acknowledgements

Anonymised for review.

References

- [1] Uhlig, H. (2022). A Luna-tic stablecoin crash. *NBER Working Paper 30256*.

- [2] Griffin, J. M. and Shams, A. (2020). Is Bitcoin really untethered? *Journal of Finance*, 75(4):1913–1964.
- [3] Hautsch, N., Scheuch, C., and Voigt, S. (2018). Building trust takes time: Limits to arbitrage for blockchain-based assets. *arXiv:1812.00595*.
- [4] Shleifer, A. and Vishny, R. W. (1997). The limits of arbitrage. *Journal of Finance*, 52(1):35–55.
- [5] Gromb, D. and Vayanos, D. (2010). Limits of arbitrage: The state of the theory. *Annual Review of Financial Economics*, 2:251–275.
- [6] Cont, R. and Kukanov, A. (2013). Optimal order placement in limit order markets. *Quantitative Finance*, 17(1):21–39.
- [7] Makarov, I. and Schoar, A. (2020). Trading and arbitrage in cryptocurrency markets. *Journal of Financial Economics*, 135(2):293–319.
- [8] López de Prado, M. (2018). *Advances in Financial Machine Learning*. Wiley.
- [9] Cintra, R. and Holloway, C. (2023). Detecting depegs: Towards safer passive liquidity provision on Curve Finance. *arXiv:2306.10612*.
- [10] Vidal-Tomás, D., Briola, A., and Aste, T. (2023). FTX’s downfall and Binance’s consolidation. *arXiv:2302.11371*.
- [11] Gatev, E., Goetzmann, W. N., and Rouwenhorst, K. G. (2006). Pairs trading: Performance of a relative-value arbitrage rule. *Review of Financial Studies*, 19(3):797–827.
- [12] Tibshirani, R. (1996). Regression shrinkage and selection via the lasso. *JRSS-B*, 58(1):267–288.
- [13] Breiman, L. (2001). Random forests. *Machine Learning*, 45(1):5–32.
- [14] Chen, T. and Guestrin, C. (2016). XGBoost: A scalable tree boosting system. In *KDD*, pp. 785–794.
- [15] Adams, R. P. and MacKay, D. J. C. (2007). Bayesian online changepoint detection. *arXiv:0710.3742*.
- [16] Moallemi, C. C. and Wang, M. (2022). A reinforcement learning approach to optimal execution. *Quantitative Finance*, 22(6):1051–1069.
- [17] Ning, B., Lin, F. H. T., and Jaimungal, S. (2021). Double deep Q-learning for optimal execution. *Applied Mathematical Finance*, 28(4):361–380.
- [18] Gogol, K., Messias, J., Miori, D., Tessone, C., and Livshits, B. (2024). Quantifying arbitrage in automated market makers: An empirical study of Ethereum ZK rollups. In *MARBLE 2024*.
- [19] Adams, A., Chan, B. Y., Markovich, S., and Wan, X. (2023). Don’t let MEV slip: The costs of swapping on the Uniswap protocol. *arXiv:2309.13648*.
- [20] Wu, W. and Liu, C. (2026). Stability anchors and risk amplifiers: Tail spillovers across stablecoin designs. *arXiv:2602.18820*.
- [21] Hernandez Cruz, W., Xu, J., Tasca, P., and Campajola, C. (2024). No questions asked: Effects of transparency on stablecoin liquidity during the collapse of Silicon Valley Bank. *arXiv:2407.11716*.
- [22] Mohl, V., Frey, S., Leyland, R., Li, K., Nigmatulin, G., Cucuringu, M., Zohren, S., Foerster, J., and Calinescu, A. (2025). JaxMARL-HFT: GPU-accelerated large-scale multi-agent reinforcement learning for high-frequency trading. In *ICAFIF '25*. *arXiv:2511.02136*.
- [23] Olby, O., Bacalum, A., Baggott, R., and Stillman, N. (2025). Right place, right time: Market simulation-based RL for execution optimisation. In *ICAFIF '25*. *arXiv:2510.22206*.