

Stablecoin StressBench: Measuring the 12× Gap Between Visible and Executable Arbitrage

Anonymous Author(s)

Abstract

Price charts show twelve times more profitable stablecoin arbitrage than a trader can capture. No prior benchmark provides per-minute executable labels, route-level depth provenance, or a hindsight oracle ceiling, so every published model score measures performance against windows that cannot be traded. We introduce Stablecoin StressBench: an 18-event source-tracked catalogue across seven mechanism classes (2020–2023), tiered data governance, per-minute VWAP book-walking labels with route provenance, and a hindsight oracle ceiling. The 12× optical-to-executable gap persists across all cost parameters and widens to near-zero at institutional scale. Across 25+ model configurations every calm-trained approach fails economically, confirming the gap is not closable by architecture or feature choices alone. Cross-mechanism meta-labeling trained on Terra/LUNA recovers +82.5 bps on the SVB test split (51.0% oracle capture), and pooling four events across three mechanism classes reaches 51.8%. A conditioned RL agent at identical positive-label density earns −29.2 bps, consistent with explicit binary supervision being the missing ingredient. All four training conditions saturate near 50% oracle capture, pointing to a structural late-stage ambiguity that additional mechanism diversity does not resolve. We release the catalogue, labels, oracle, and model harness under CC-BY 4.0.

Keywords

stablecoin arbitrage, market microstructure, execution-aware labels, financial benchmarks, cryptocurrency

ACM Reference Format:

Anonymous Author(s). 2026. Stablecoin StressBench: Measuring the 12× Gap Between Visible and Executable Arbitrage. In *Proceedings of ACM International Conference on AI in Finance (ICAIF '26)*. ACM, New York, NY, USA, 8 pages.

1 Introduction

Stablecoin stress is not one thing. From 2020 to 2023 we identify 18 distinct episodes spanning seven mechanism classes: algorithmic death spirals (Terra/LUNA, IRON/TITAN), fiat-reserve shocks (USDC/SVB), regulatory winddowns (BUSD), exchange-credit collapses (FTX, Celsius), DeFi pool imbalances (Curve/USDT), collateral liquidations (DAI Black Thursday), and RWA failures (USDR, Acala). Every prior stablecoin stress study examines at most one of these mechanisms and none provide per-minute executable labels.

The paper follows one spine. Price charts show profitable arbitrage that vanishes under execution (Gap 1, 12×). The surviving executable windows cannot be found by any calm-trained model (Gap 2). Only cross-mechanism supervision bridges Gap 2, and even then performance saturates near 50% of what a hindsight oracle achieves. Figure 2 shows all three layers for the SVB test window

and traces one representative false-positive minute from apparent profit to net loss.

We call dislocations visible on charts but absent in execution **optical**, and surviving windows **executable**. The 12× ratio is measured at the 10 bps threshold and \$10K notional; it rises to 14× under a 20% depth haircut and collapses to near-zero at \$500K, so the headline number understates the constraint for institutional participants.

The execution-grade analysis centers on one event: the USDC/SVB stress window of March 2023. The 18-event catalogue situates that event in a broader stress universe and motivates cross-mechanism training by showing that depth-withdrawal and spread-widening recur across mechanistically distinct triggers.

Contributions.

- (1) **First execution-aware stablecoin benchmark.** An 18-event source-tracked catalogue across seven mechanism classes, tiered data governance (Tier A: VWAP labels and oracle for USDC/SVB; Tier B: price/liquidity; Tier C: taxonomy), per-minute VWAP book-walking labels with route provenance, and a hindsight oracle ceiling. (§3–5)
- (2) **12× gap and cross-mechanism transfer.** Gap 1: 34.3% optical vs. 2.88% executable at \$10K. Gap 2: every calm-trained model fails; meta-labeling trained on Terra/LUNA achieves +82.5 bps on SVB (51.0% oracle capture) and pooling four events reaches 51.8%. All four conditions saturate near 50%, pointing to a structural late-stage ambiguity that additional mechanism pooling cannot lift. (§6)
- (3) **Supervision format, not label density, as binding constraint.** A conditioned PPO-GRU at 17% positive density earns −29.2 bps vs. meta-labeling’s +82.5 bps at identical density, consistent with explicit binary labels being the ingredient reward optimisation cannot replicate; we treat this as a hypothesis for the multi-mechanism RL experiment in future work. (§6.3.2)

2 Related Work

Stablecoin stress and depeg detection. Uhlig [1] and Griffin and Shams [2] analyse Terra/LUNA and stablecoin non-fundamental demand. Hernandez Cruz et al. [21] study DEX liquidity during the SVB collapse without modelling CEX execution costs. Cintra and Holloway [9] detect the Terra depeg 12 hours early via BOCPD on AMM reserves; we replicate their approach on CEX cross-quote data and find AUROC 0.229, confirming that AMM and CEX signals differ fundamentally. Wu and Liu [20] study tail spillovers across stablecoin designs without per-minute execution analysis. None of these works provide VWAP-grade execution labels, a hindsight oracle, or cross-mechanism transfer evaluation.

Limits to arbitrage and execution frictions. Shleifer and Vishny [4], Gromb and Vayanos [5], and Makarov and Schoar [7] show that visible price gaps need not be profitable once execution

Table 1: Route-leg depth provenance.

Leg	SVB Mar 2023	2024 windows
Buy BTC-USDC	kline-proxy [†]	real L2 (futures)
Cross USDC-USDT	partial/proxy; real L2	real L2 (futures)
	where archived	
Sell BTC-USDT	raw/futures depth	real L2 (futures)
Ref BTC-USD	candle-proxy	candle-proxy

perpetual listed Jan 2024; all 2022 windows kline-proxy.

Table 2: Dataset splits (56,134 one-minute bars).

Split	Event	Min.	Pos.
Train	Calm control (×3)	28,776	0.00%
Validation	Terra/LUNA May 2022	11,526	2.30%
Test	USDC/SVB Mar 2023	15,832	2.88%

frictions are included. Our 12× ratio is this effect measured at one-minute resolution. Hautsch et al. [3] quantify settlement latency as a first-order cost on blockchain assets. Cont and Kukanov [6] establish VWAP book-walking as the correct profitability measure. Gogol et al. [18] measure Maximal Arbitrage Value on AMM-to-CEX routes; their partial execution labels cover AMM venues only and lack a hindsight oracle. Adams et al. [19] study swap costs on Uniswap without stress-event execution analysis.

ML and RL for execution. Lopez de Prado [8] introduces meta-labeling, which we extend to the cross-mechanism transfer setting. Moallemi and Wang [16] and Ning et al. [17] supply the finite-horizon MDP framing we adopt. Mohl et al. [22] and Olby et al. [23] operate in reactive LOB environments with market-impact feedback; our diagnostic uses historical replay without feedback, isolating training supervision as the variable.

3 Benchmark Design

3.1 Data Sources and Instruments

The 18-event catalogue defines the historical stress universe; the six-window empirical panel (56,134 one-minute bars) defines the model dataset; and the Tier-A SVB windows define the execution-grade testbed. Price features are sourced from Binance Vision (klines and aggTrades) and Coinbase REST candles. Net-profit labels use three depth legs: Binance BTC-USDC (buy), Binance BTC-USDT (sell), and Coinbase BTC-USD (reference). The BTC-USDC buy leg uses kline-proxy depth for the SVB window because the perpetual contract did not exist on Binance until January 2024; a 20% depth haircut confirms robustness (gap widens to 14×, oracle stays above +130 bps).

Data quality varies by leg; every record carries a provenance tag (Table 1).

3.2 Event-Based Split Design

Splits are defined by event identity, so no model trained on the calm split has encountered stress dynamics. Threshold calibration uses the Terra/LUNA validation split, an independent stress event, so the SVB test split is never observed during model development. The 2.88% positive rate in the test split implies severe class imbalance. We therefore report economic net bps as the primary criterion; AUROC and AUPRC are secondary.

Five leakage checks confirm validity: label shuffles produce zero or negative P&L; one-minute lookahead does not improve performance; the calm training split contains zero executable windows; train and test are separated by event identity; and cross-mechanism calibration uses only the Terra/LUNA calibration block.

3.3 Historical Catalogue: Situating SVB in the Stress Universe

The execution-grade analysis centers on USDC/SVB. The 18-event catalogue motivates this choice and justifies cross-mechanism training: depth withdrawal and spread widening recur across all mechanism classes, so a model that has seen them in one event can recognise them in another.

Table 3 lists all 18 events with source-verified statistics (“est.” denotes values from verified secondary sources, not per-minute data). Tier-A events support VWAP execution labels and oracle bounds; Tier-B support price/liquidity claims; Tier-C contribute taxonomy coverage. The pooled transfer experiment draws from Terra/LUNA (algorithmic), Celsius/3AC and FTX (exchange credit), and BUSD (regulatory), all held mechanistically distinct from the Tier-A SVB test event.

Figure 1 shows the event universe as a bubble timeline: mechanism class on the vertical axis, date on the horizontal, and bubble area proportional to the absolute maximum depeg.




Figure 1: 18-event stress universe. Each bubble is one event; area $\propto$ |max depeg|; gold edge indicates above-peg stress. Tier A (navy); Tier B (blue); Tier C (grey). Basis fingerprints for Terra/LUNA and USDC/SVB appear in Figure 9.

3.4 Benchmark Positioning

Table 4 positions StressBench against closest prior work. No prior work provides per-minute executable CEX labels, route-level L2 provenance, a hindsight oracle ceiling, and cross-mechanism transfer evaluation together.

Table 3: StressBench event catalogue: all 18 events, chronological. Max depeg in bps (+ = above peg; † = percentage-point collapse to zero). “est.” = source-verified estimate, not per-minute data. Tier A = execution-grade (VWAP labels + oracle); Tier B = price/liquidity claims; Tier C = taxonomy only.

Event	Date	Coin	Mechanism	Max dep.	Dur.	T	Event	Date	Coin	Mechanism	Max dep.	Dur.
DAI Black Thursday	Mar '20	DAI	Collateral	est. +150	days	B	Acala aUSD Exploit	Aug '22	aUSD	RWA/Niche	est. -99pp [†]	days
FEI Launch Stress	Apr '21	FEI	Algo./Reflex.	est. -20pp [†]	days	C	HUSD Issuer Failure	Aug '22	HUSD	Exch. Credit	est. -800	days
IRON/TITAN Collapse	Jun '21	IRON	Algo./Reflex.	est. -100pp [†]	hrs	C	Binance USDC → BUSD Conv.	Sep '22	BUSD	Regulatory	est. +0	days
MIM/Wonderland Shock	Jan '22	MIM	Algo./Reflex.	est. -300	days	C	FTX Collapse	Nov '22	USDT	Exch. Credit	est. -50	days
Terra/UST Collapse	May '22	UST	Algo./Reflex.	est. -100pp [†]	wks	B	BUSD Regulatory Winddown	Feb '23	BUSD	Regulatory	est. -30	days
Curve 3Pool/UST Stress	May '22	UST	DeFi Pool	est. -500	days	B	USDC/DAI DeFi Secondary	Mar '23	USDC	DeFi Pool	est. -200	days
Celsius/3AC Contagion	Jun '22	USDT	Exch. Credit	est. -100	wks	B	USDC/SVB Stress	Mar '23	USDC	Fiat Reserve	-13pp [†]	days
USDD/TRON Stress	Jun '22	USDD	Algo./Reflex.	est. -200	days	B	USDC Recovery (Post-SVB)	Mar '23	USDC	Fiat Reserve	-10	days
							USDT/Curve Pool Stress	Jun '23	USDT	DeFi Pool	est. -80	days
							USDR RWA Failure	Oct '23	USDR	RWA/Niche	est. -50pp [†]	days

2A, 11B, 5C events. Tier-A VWAP labels available only for USDC/SVB (primary benchmark) and its recovery window.

Table 4: Comparison with related benchmarks. ✓ full; ~ partial; — absent.

Work	Stress	Exec. labels	L2 prov.	Oracle	Cross-event
Cintra & Holloway [9]	✓	—	AMM	—	✓
Hernandez Cruz et al. [21]	✓	—	DEX	—	—
Gogol et al. [18]	—	~	CEX	MAV	—
Adams et al. [19]	—	~	CEX	—	—
StressBench (ours)	✓	✓	CEX	✓	✓

3.5 Feature Sets

Four nested sets isolate marginal signal value. `price_only` provides 5 cross-quote basis columns. `price_plus_book` adds 7 microstructure signals (spread, depth, imbalance). `price_book_frag` adds 3 cross-venue fragmentation metrics. `price_book_settle` adds 7 on-chain settlement proxies, which are kline-proxy approximations for the SVB and Terra/LUNA windows (Tier B).

4 Execution-Aware Labels

4.1 Three-Layer Framework

Figure 2 structures the paper’s central argument visually. Layer 1 (optical) contains every minute where the price chart shows a profitable gap. Layer 2 (executable) contains the minutes that remain profitable after VWAP book-walking, taker fees, and settlement latency. Layer 3 (captured) contains the minutes a deployed model can identify in advance. Gap 1 is the distance from Layer 1 to Layer 2; Gap 2 is the distance from Layer 2 to Layer 3. Both gaps are nonzero, independent, and large. Panel (b) of Figure 2 traces one representative false-positive minute through the full cost stack, showing why 34.3% optical activity collapses to 2.88% executable.

4.2 Cross-Quote Basis and Labels

Three equations define the label construction.

$$b_{\text{USDC}}(t) = 10,000 \times \frac{P_{\text{BTC/USDC}}(t) - P_{\text{BTC/USD}}(t)}{P_{\text{BTC/USD}}(t)} \quad [\text{bps}] \quad (1)$$

$$b_{\text{max}}(t) = \max(|b_{\text{USDC}}(t)|, |b_{\text{USDT}}(t)|) \quad (2)$$

$$y_{\text{basis}}(t) = 1[|b_{\text{USDC}}(t+h)| > \tau] \quad (3)$$

Primary task: `basis_usdc_1m_gt10bps` ($\tau=10$ bps, $h=1$ min).




Figure 2: (a) Three-layer framework for the SVB test window (15,832 min). Layer 1 optical (34.3%), Layer 2 executable (2.88%), Layer 3 captured by cross-mechanism meta-labeling (1.47%). Gap 1 (12×) is a measurement problem; Gap 2 is a training problem. (b) Worked example: a minute that appears profitable on the chart nets -15.8 bps after VWAP book-walking, fees, and route-direction mismatch.




Figure 3: Cross-quote basis during the SVB test window (Mar 10–20 2023). USDC basis (navy): 12.4% of minutes above 10 bps; USDT basis (red): 27.4%; max-abs (gold): 34.3%. The price-to-execution ratio is 12×.

4.3 VWAP Execution and Net Profit Labels

$$\text{net}(q, t) = 10,000 \times \left[\frac{\text{VWAP}_{\text{sell}} - \text{VWAP}_{\text{buy}}}{P_{\text{ref}}} - f_{\text{buy}} - f_{\text{sell}} - \delta_{\text{settle}} \right] \quad [\text{bps}] \quad (4)$$

$$y_{\text{arb}}(q, h, t) = 1 \left[\max_{s \in [t, t+h]} \text{net}(q, s) > c \right] \quad (5)$$

Profit threshold $c=10$ bps, taker fee 4 bps per side, and $\delta=5$ bps. The δ assumption is validated against 100K on-chain USDC ERC-20 transfers: median cost was 1.4–3.2 bps on nine of ten SVB days; the

one exception (Mar 11 gas spike) saw P95 reach 8.6 bps.

$$\text{Oracle}(q) = \frac{1}{|\mathcal{T}^+|} \sum_{t \in \mathcal{T}^+} \text{net}(q, t), \quad \mathcal{T}^+ = \{t : \text{net}(q, t) > c\} \quad (6)$$

Oracle capture $\kappa_M = \bar{r}_M / \text{Oracle}(q)$ is the fraction of the oracle’s per-trade return recovered by model M .

5 Models and Evaluation

5.1 Model Ladder

The ladder is a sequential elimination experiment. Before applying the cross-mechanism fix, six remedy families are tested in turn. (1) *Architecture*. GRU and MLP-Window fail despite achieving the highest AUROC among learned sequence models. (2) *Feature engineering*. All four nested feature sets produce negative returns; adding depth and settlement signals does not help when the training split contains zero stress-regime examples. (3) *Class balancing*. Cost-sensitive reweighting cannot synthesise stress dynamics from calm data. (4) *Regression targeting*. ExpNetProfitRegressor directly optimises the economic criterion and still reaches -73.8 bps, ruling out label mismatch. (5) *Regime detection*. BOCPD, CUSUM, and EWMA all fail because CEX basis is mean-reverting (BOCPD AUROC 0.229). (6) *RL reward optimisation*. A conditioned PPO-GRU at identical positive-label density earns -29.2 bps versus meta-labeling’s $+82.5$ bps. Every family fails for the same structural reason: distributional mismatch between calm training and stress test.

The cross-mechanism extension of meta-labeling operates in three steps.

- (1) **Primary filter**. Flag t if $|b_{\text{USDC}}(t)| > \tau_1$ ($\tau_1=10$ bps).
- (2) **Secondary classifier**. Fit f_θ on $\{(x_t, y_{\text{arb}}(t))\}_{t \in \mathcal{E}_{\text{train}}}$, where $\mathcal{E}_{\text{train}}$ is a mechanistically distinct stress event.
- (3) **Entry**. Signal at t iff primary fires and $f_\theta(x_t) > \theta^*$, where $\theta^* = \arg \max_\theta \sum_{i: \hat{p}_i > \theta} \text{net}(q, t_i)$ on the Terra/LUNA calibration block.

5.2 RL Execution Policy Baseline

The entry-timing problem is framed as a finite-horizon MDP. At each minute t the agent observes the 30-step lookback of price_plus_book features, selects $a_t \in \{\text{enter}, \text{wait}\}$, and receives reward $r_t = \text{net}(q, t)$ on entry, 0 on wait. We train a PPO agent with a GRU policy network (64 units, 30-step lookback) using the cross-mechanism protocol: training on Terra/LUNA stress windows only and evaluating on USDC/SVB without retraining. This directly tests whether the sequential timing framing overcomes the distributional mismatch that defeats all calm-trained models.

5.3 Calibration and Evaluation

Net bps denotes mean net basis points per triggered trade after VWAP, taker fees, and settlement penalties. Thresholds are calibrated on the validation split by maximising $\theta^* = \arg \max_\theta \sum_{i: \hat{p}_i > \theta} \text{net}(q, t_i)$ subject to a minimum of 25 trades. Primary metrics are mean net bps per triggered trade and oracle capture κ_M . Secondary metrics are AUROC and AUPRC. NoTrade at 0 bps is the economic floor.

Positive density clarification. Three densities appear in the paper. 2.88% is the fraction of SVB test minutes that are executable.

Table 5: Master oracle and gap table. Every oracle value cited in the paper appears here. Gap = oracle minus best deployable model.

Task / Window	N	Oracle	Best model	Gap	AUPRC
basis_usdc_1m (full test)	15,832	+161.7 bps	0 bps ^a	162 bps	0.253
exec_arb_q10k_5m (full test)	15,832	+224.6 bps	-41.1 bps	266 bps	0.060
exec_arb_q50k_5m (full test)	15,832	+146.8 bps	-112.7 bps	260 bps	0.063
SVB stress phase only (Mar 10–14)	7,189	+259.9 bps	—	—	—
Terra/LUNA validation split	11,526	+87.6 bps	—	—	—

^aBest calibrated model is NoTrade (0 trades).

Kline-proxy buy leg; 20% haircut bounds in §3.1.

Table 6: Price-to-execution gap (\$10K, 1-min horizon, SVB test split).

Thresh.	P _{max}	P _{USDC}	Exec.	Ratio
0 bps	95.86%	82.96%	3.34%	29×
5 bps	61.99%	25.32%	3.03%	20×
10 bps	34.33%	12.45%	2.88%	12×
25 bps	9.54%	6.48%	2.62%	3.6×

2.30% is the fraction of the full Terra/LUNA stream that is executable. 17% is the fraction of primary-signal minutes (those where $|b_{\text{USDC}}| > 10$ bps) that are executable in Terra/LUNA; this is the density seen by the conditioned PPO-GRU and the meta-labeler secondary classifier during training. The RL-vs-meta-labeling comparison holds density fixed at the 17% figure.

Five evaluation tasks ship with StressBench v1.0 (train: calm or Terra block; val: Terra calibration; test: SVB): optical_basis (AUROC/AUPRC), exec_arb_q10k and exec_arb_q50k (net bps, oracle capture), cross_mech (net bps), and policy_entry (net bps, trades). Valid submissions must report net bps per trade, oracle capture, cost assumptions, notional size, and data tier used. Models are ranked by net bps, not AUROC.

6 Results

Table 5 maps every oracle figure in the paper to its task and window. All subsequent references point back to one of these rows.

6.1 The 12× Execution Barrier

At a 10 bps threshold, 34.3% of test minutes are optically dislocated but only 2.88% are executable after VWAP slippage, taker fees, and settlement latency. The ratio persists across all basis levels and widens with notional size: the executable rate falls from 2.88% at \$10K to 0.03% at \$500K (Figure 4). Book depth, not fees, prices out institutional participants.

Block-bootstrap CIs (60-min blocks, 2,000 resamples): executable rate 95% CI [1.47, 4.59]%; ratio [7.6, 23.0]×; oracle [+213.9, +304.2] bps.

Table 7 shows the stress/recovery split. All executable windows fall in the SVB stress phase; the recovery phase contains zero executable windows despite 21.6% optical activity, confirming that depth returns before the basis closes. Terra/LUNA validation independently replicates the gap at 5.9×, providing directional evidence that the finding is not SVB-specific.

The gap is robust across the full cost-parameter space: sweeping taker fees from 2 to 10 bps and settlement penalty from 0 to 10 bps moves the executable rate between 2.84% and 3.66%.

Table 7: Optical and executable diagnostics by window. SVB stress/recovery are Tier-A execution-grade. Terra/LUNA is validation-grade only.

Window	Min.	Optical	Exec.	Oracle
SVB Stress (Mar 10–14)	7,189	53.0%	6.63%	+259.9 bps
SVB Recovery (Mar 15–20)	8,643	21.6%	0.00%	—
Terra/LUNA (May 2022, val.)	11,526	14.0%	2.40%	+87.6 bps




Figure 4: Notional scaling (SVB test split). The executable rate collapses from 2.88% at \$10K to 0.03% at \$500K: book depth, not fees, is the binding institutional constraint.

6.2 Calm-Training Failure

The central finding is not that individual models fail, but that all six families converge to the same result: no architecture overcomes the distributional mismatch between calm training and stress test. This convergence across 25+ configurations is itself the diagnostic. Table 8 shows the tabular model results on `basis_usdc_1m_gt10bps`; oracle gap values map to Table 5.

Price-threshold rules overtrade massively (607–1,969 trades at –136 to –347 bps). Route-direction mismatch is the primary failure mode: 581 of 1,851 false positives are triggered by USDT dislocations where the USDC route is simultaneously non-executable at –38.7 bps each. True-positive windows carry a mean USDC basis of 344 bps; false positives carry just 0.56 bps and show higher average bid depth (72,805 vs. 59,961 BTC units), ruling out thin books as the cause.

The discrimination–profitability disconnect. Among sequence models, the GRU achieves the highest AUROC (0.80) yet earns the worst economic return (–239 bps). Figure 5 shows that this inversion holds across every model family: high discriminative accuracy does not imply profitability under 2.88% class imbalance. The MLP-Window reaches –341 bps.

Figure 7 decomposes the 260 bps oracle gap (Table 5) into three measurable components. False-positive exposure is the dominant drag: route-direction mismatch accounts for 581 of 1,851 false positives from price-rule trading. The remaining gap splits into timing error (40–80 bps, from entering 1–2 minutes late in a declining depth regime) and sizing error (40–100 bps, as depth depletion shrinks the executable fraction from 2.88% at \$10K to 1.62% at \$50K). Meta-labeling’s 51% hit rate versus 6% for the price rule substantially reduces route-direction mismatch, though a 49% false-positive rate and roughly 113 missed profitable windows confine net returns to




Figure 5: AUROC vs. net bps per triggered trade (SVB test split, all families). High discriminative accuracy does not imply profitability: GRU achieves AUROC 0.80 yet –239 bps; oracle trades at AUROC 0.52 and +162 bps. Only meta-labeling (diamond) crosses zero.

Table 8: Tabular model ladder (`basis_usdc_1m_gt10bps`). Hit% is fraction of executable trades. †Oracle ceiling. *0-trade degenerate.

Model	Feat.	Net bps	Trades	AUROC	Hit%
Oracle [†]	—	+161.7	316	—	100%
lgbm*	all	—	0	0.653	—
logistic	all	–49.1	1	0.133	0%
gross_arb	px	–136.0	611	0.531	4%
price_10	px	–269.5	2002	0.833	6%
price_25	px	–347.3	1025	0.746	5%
no_trade	—	0.0	0	—	—
ExpNet_lgbm	px	–73.8	537	—	6%
ExpNet_xgb	px	–57.2	70	—	9%
lgbm_weighted	pb	–65.9	2	0.497	0%
xgb_weighted	px	–90.5	30	0.631	0%
rf_balanced	px	–89.2	11	0.467	0%

px = price_only; pb = price_plus_book; all = price_book_settle.

+82.5 bps. The prediction gap is not a model-selection problem. It is a training-distribution problem, and the only empirically validated solution is cross-mechanism supervision.

Table 9: Failure mode diagnosis by model family. All families fail for the same root cause: distributional mismatch.

Family	Failure mode	Key evidence
Price rules	Route mismatch	581/1,851 FP: USDT route dislocated, USDC at –38.7 bps
Tabular ML	Degenerate threshold	LightGBM: 0 trades; no threshold yields ≥25 with positive return
Sequence (GRU)	AUROC–P&L inversion	AUROC 0.80 (best sequence), net –239 bps (worst); 401 bps gap
RL (PPO-GRU)	Timing mismatch	919 trades at –29.2 bps; Terra pattern fires on wrong SVB windows
Regime detection	Mean-reversion violation	BOCPD AUROC 0.229; CEX basis reverts, not shifts

6.3 Cross-Mechanism Transfer

Cross-mechanism meta-labeling trained on Terra/LUNA is the only method that crosses zero, earning +82.5 bps (51.0% oracle capture




Figure 6: No calm-trained model ever crosses zero. Meta-labeling trained on Terra/LUNA is the only approach that achieves positive returns, earning +82.5 bps against the oracle’s +161.7 bps. All other learned models lose money despite sometimes high AUROC. Oracle (gold); meta-label Terra/LUNA (green solid); meta-label 4-event pooled (green dotted); PriceBasis10bps (red); NoTrade (grey dashed). Shaded region: SVB shock window.




Figure 7: Oracle gap decomposition (\$10K notional, SVB test split). False-positive drag is directly measured from false-positive trades. Timing error (40–80 bps) and sizing error (40–100 bps) are estimated from entry-delay and notional-scaling analysis.

on the `basis_usdc_1m` task, oracle +161.7 bps from Table 5). A reinforcement learning agent given the same 17% primary-signal density earns −29.2 bps, consistent with explicit binary supervision being the binding ingredient; we treat this as a hypothesis rather than a causal claim (Section 6.3.2).

6.3.1 Cross-Mechanism Transfer and Training Diversity. Meta-labeling achieves positive returns on every non-degenerate training condition tested (Table 10). Four-event pooled training spans algorithmic, exchange-credit, and regulatory mechanisms and reaches +83.7 bps (51.8% oracle capture).

Cross-mechanism training succeeds because depth withdrawal, spread widening, and route mismatch recur across stress triggers even when triggers differ. Figure 9 shows this directly: ask-side depth collapses and spread widens identically in primary-signal windows across Terra/LUNA and SVB despite mechanistically distinct causes. The order book responds to uncertainty, not its source.

Table 10: Transfer and policy results (`basis_usdc_1m_gt10bps`, `price_plus_book`, SVB test split). Mean net bps per triggered trade; single-event block calibration throughout. DGP robustness ranges (500-seed bootstrap): Terra [78.7, 86.2], Celsius [74.1, 89.9], Pooled [75.3, 93.1].

Training source	Mechanism	Net bps	Trades	Oracle cap.
Oracle	hindsight	+161.7	316	100%
Calm-control	none	0.0	0	0%
Terra/LUNA only	algorithmic	+82.5	397	51.0%
Celsius/3AC only	exch. credit	+79.7	221	49.1%
FTX only*	exch. credit	+36.5	1	22.5%
Four-event pooled	alg.+exch.+reg.	+83.7	163	51.8%
PPO-GRU conditioned	RL on Terra	−29.2	919	−18%

*FTX: signal intensity < 5 bps; mechanism-intensity failure.

SHAP analysis (Table 11 and Figure 10) confirms the two-stage structure: price-basis features supply the candidate signal (94.8% of importance), while ask-side depth and spread filter execution quality (3.4%) once the model has seen stress-regime depth withdrawal.

Table 11: Top SHAP importances, meta-labeling secondary classifier (LightGBM, Terra to SVB). Basis dominates; book features filter execution quality within primary-signal windows.

Feature	Importance	Group
USDC cross-quote basis	49.6%	price
Primary basis	32.8%	price
Deviation from \$1	10.1%	price
USDT cross-quote basis	2.0%	price
Ask depth (10 bp)	0.5%	book/frag
Bid-ask spread	0.5%	book/frag
Price/basis total	94.8%	
Book/fragmentation total	3.4%	
Settlement total	1.8%	

Three regime-detection filters (CUSUM, EWMA, BOCPD [15]) were evaluated as alternative pre-filters. BOCPD achieves AUROC 0.229 on CEX cross-quote data. CEX basis is mean-reverting rather than change-point structured, so changepoint methods cannot substitute for the training-distribution fix that meta-labeling provides.

The 50% saturation. All four training conditions produce oracle capture in the range 51.0–51.8% (Figure 8). Pooling four mechanistically distinct events does not lift this range. In the minutes immediately before a dislocation closes, ask-side depth partially recovers and the basis compresses, making late-stage windows microstructurally indistinguishable from genuine opportunities. The secondary classifier cannot resolve this ambiguity regardless of training diversity. We report this as an observed saturation across our four conditions, not as a proven ceiling; bootstrapping the four capture rates produces overlapping 95% intervals, consistent with a true value near 50% but not ruling out higher capture with substantially more training diversity. Closing the gap further requires early-entry timing labels or a reactive simulator that models the agent’s own market impact.

6.3.2 RL Policy Failure. This is a controlled comparison with one varying factor: how the profitable-window signal is delivered. Meta-labeling receives explicit binary labels ($y=1$ for each executable




Figure 8: Oracle capture across four training conditions. All saturate near 50% (dashed). FTX (hatched) is near-degenerate.




Figure 9: Ask depth (blue) collapses and spread widens (red) identically in Terra/LUNA and SVB. The order book responds to uncertainty, not its cause.




Figure 10: SHAP importances (Table 11). Basis features dominate (94.8%).

window); the conditioned PPO-GRU receives a sparse reward signal over the same windows. Both see identical training data at the same 17% primary-signal density. Despite this, the two methods differ by 111.7 bps in net return. This difference is consistent with explicit binary supervision being the binding ingredient that reward optimisation cannot replicate at sparse positive rates. We treat this as a hypothesis rather than a causal claim; a definitive test requires multiple RL runs across the pooled multi-mechanism protocol, which we leave as the most immediate next experiment.

Without conditioning, the PPO-GRU degenerates to zero entries: the 2.3% positive rate in the full Terra/LUNA stream is too sparse to calibrate the reward signal. Conditioning on primary-signal windows resolves the density problem, but the agent still earns −29.2 bps on 919 trades. The 919-trade overtrading pattern is consistent with a timing-transfer mismatch: the agent learned Terra/LUNA’s depth-withdrawal sequence and fires on it in SVB even when the timing does not match.

7 Limitations

Single Tier-A event. Execution-grade VWAP labels exist only for USDC/SVB and its recovery window. Extending to exchange-credit or DeFi events requires route-complete L2 archives not yet available.

Kline-proxy buy leg. Covered in §3.1: 20% haircut widens the gap to 14×, confirming robustness.

RL market-impact feedback. The conditioned PPO-GRU’s −29.2 bps may reflect a timing-transfer mismatch (Terra/LUNA depth-withdrawal sequence firing at the wrong moments in SVB) rather than a fundamental ceiling. Pooled multi-mechanism training on the protocol established here is the direct experiment to test whether this mismatch explains the gap. Training on historical replay without market-impact feedback also miscalibrates sizing in thin-book regimes, because the agent sees depth as it was rather than as it would be after its own orders clear liquidity. Integrating the cross-mechanism initialisation protocol into a reactive simulator resolves both issues simultaneously and is the natural next step toward closing the 111.7 bps gap between meta-labeling and the conditioned RL agent.

Tier-A coverage. Execution-grade VWAP labels currently exist only for USDC/SVB and its recovery window. Extending to

exchange-credit or DeFi events requires route-complete L2 archives not yet available; this is the primary bottleneck for expanding the benchmark’s execution-grade scope.

Release constraints and data ethics. The benchmark releases derived features, labels, event windows, source-tracking meta-data, scripts, and frozen paper outputs under CC-BY 4.0. Some raw exchange order-book archives are API-limited, delayed, or subscription-gated; StressBench records route provenance and depth-source tags so future users can distinguish reproducible derived claims from raw-data acquisition constraints. No private user-level trading data is included.

8 Conclusion

Stablecoin arbitrage benchmarks built on price labels measure the wrong object. The 12× optical-to-executable gap means that every published AUROC or F1 score on basis data overstates deployable performance by a factor of twelve or more. Gap 1 is not a model problem; it is a labeling problem. Gap 2 is not a feature engineering problem; it is a training distribution problem. The only empirically validated solution to Gap 2 is cross-mechanism supervision, and that solution saturates near 50% oracle capture under current conditions.

What this implies for the field: improving on StressBench requires either expanding Tier-A coverage to additional events (so that transfer can be tested, not just assumed) or demonstrating that a reactive RL agent with market-impact feedback can close the gap that historical replay cannot. The oracle ceiling provides a principled performance target for both directions. We release the catalogue, labels, oracle, and harness so that any model claiming to capture stablecoin stress arbitrage can be evaluated against the correct denominator.

Acknowledgements

Anonymised for review.

References

- [1] Uhlig, H. (2022). A Luna-tic stablecoin crash. *NBER Working Paper 30256*.
- [2] Griffin, J. M. and Shams, A. (2020). Is Bitcoin really untethered? *Journal of Finance*, 75(4):1913–1964.
- [3] Hautsch, N., Scheuch, C., and Voigt, S. (2018). Building trust takes time: Limits to arbitrage for blockchain-based assets. *arXiv:1812.00595*.
- [4] Shleifer, A. and Vishny, R. W. (1997). The limits of arbitrage. *Journal of Finance*, 52(1):35–55.
- [5] Gromb, D. and Vayanos, D. (2010). Limits of arbitrage: The state of the theory. *Annual Review of Financial Economics*, 2:251–275.
- [6] Cont, R. and Kukanov, A. (2013). Optimal order placement in limit order markets. *Quantitative Finance*, 17(1):21–39.
- [7] Makarov, I. and Schoar, A. (2020). Trading and arbitrage in cryptocurrency markets. *Journal of Financial Economics*, 135(2):293–319.
- [8] López de Prado, M. (2018). *Advances in Financial Machine Learning*. Wiley.
- [9] Cintra, R. and Holloway, C. (2023). Detecting depegs: Towards safer passive liquidity provision on Curve Finance. *arXiv:2306.10612*.
- [10] Vidal-Tomás, D., Briola, A., and Aste, T. (2023). FTX’s downfall and Binance’s consolidation. *arXiv:2302.11371*.
- [11] Gatev, E., Goetzmann, W. N., and Rouwenhorst, K. G. (2006). Pairs trading: Performance of a relative-value arbitrage rule. *Review of Financial Studies*, 19(3):797–827.
- [12] Tibshirani, R. (1996). Regression shrinkage and selection via the lasso. *JRSS-B*, 58(1):267–288.
- [13] Breiman, L. (2001). Random forests. *Machine Learning*, 45(1):5–32.
- [14] Chen, T. and Guestrin, C. (2016). XGBoost: A scalable tree boosting system. In *KDD*, pp. 785–794.
- [15] Adams, R. P. and MacKay, D. J. C. (2007). Bayesian online changepoint detection. *arXiv:0710.3742*.

- [16] Moallemi, C. C. and Wang, M. (2022). A reinforcement learning approach to optimal execution. *Quantitative Finance*, 22(6):1051–1069.
- [17] Ning, B., Lin, F. H. T., and Jaimungal, S. (2021). Double deep Q-learning for optimal execution. *Applied Mathematical Finance*, 28(4):361–380.
- [18] Gogol, K., Messias, J., Miori, D., Tessone, C., and Livshits, B. (2024). Quantifying arbitrage in automated market makers: An empirical study of Ethereum ZK rollups. In *MARBLE 2024*.
- [19] Adams, A., Chan, B. Y., Markovich, S., and Wan, X. (2023). Don't let MEV slip: The costs of swapping on the Uniswap protocol. *arXiv:2309.13648*.
- [20] Wu, W. and Liu, C. (2026). Stability anchors and risk amplifiers: Tail spillovers across stablecoin designs. *arXiv:2602.18820*.
- [21] Hernandez Cruz, W., Xu, J., Tasca, P., and Campajola, C. (2024). No questions asked: Effects of transparency on stablecoin liquidity during the collapse of Silicon Valley Bank. *arXiv:2407.11716*.
- [22] Mohl, V., Frey, S., Leyland, R., Li, K., Nigmatulin, G., Cucuringu, M., Zohren, S., Foerster, J., and Calinescu, A. (2025). JaxMARL-HFT: GPU-accelerated large-scale multi-agent reinforcement learning for high-frequency trading. In *ICAIF '25*. *arXiv:2511.02136*.
- [23] Olby, O., Bacalum, A., Baggott, R., and Stillman, N. (2025). Right place, right time: Market simulation-based RL for execution optimisation. In *ICAIF '25*. *arXiv:2510.22206*.