

Stablecoin StressBench: An Execution-Aware Benchmark for Stablecoin Dislocations

Nigel Li

nl2992@columbia.edu

Columbia University, MAFN

New York, USA

ABSTRACT

Stablecoin stress events generate frequent price dislocations, but many apparent arbitrage opportunities disappear once executable depth, VWAP book-walking, fees, and market impact are included. We introduce Stablecoin StressBench, a transaction-cost-aware benchmark for evaluating whether AI and econometric models can distinguish optical arbitrage from executable arbitrage. During the March 2023 USDC/SVB test window, 35.1% of 1-minute windows exceed a 10bps primary cross-quote basis threshold (12.65% for the USDC-specific basis), yet only 2.88% remain profitable at \$10K notional after a full VWAP order-book walk a 12× price-to-execution gap. A hindsight oracle earns +161.7bps per trade, but every tested rule-based and ML model loses money out of sample, revealing an oracle gap exceeding 210bps. A direct expected-net-profit regressor also fails (−61.4bps), confirming the gap is structural, not a threshold artefact. The benchmark reframes stablecoin arbitrage prediction as an execution-barrier problem, and provides an open evaluation framework for future models.

KEYWORDS

stablecoin arbitrage, market microstructure, execution-aware labels, financial benchmarks, cryptocurrency

ACM Reference Format:

Nigel Li. 2026. Stablecoin StressBench: An Execution-Aware Benchmark for Stablecoin Dislocations. In *Proceedings of ACM International Conference on AI in Finance (ICAIF '26)*. ACM, New York, NY, USA, 5 pages.

1 INTRODUCTION

Stablecoin de-peg events produce dramatic price gaps visible on any chart: during the USDC/SVB crisis of March 2023, the cross-quote spread between BTC-USDC and BTC-USDT routes exceeded 100bps for sustained periods [1]. Such gaps appear to offer mechanical arbitrage opportunities, and the ML-for-finance literature has begun building models to detect and predict them [10]. The implicit assumption is that price dislocations map directly to tradeable profit.

They do not. Executable arbitrage requires clearing order-book depth at all intermediate venues, absorbing VWAP slippage as the book is walked, paying taker fees on each leg, and waiting through settlement latency before the round-trip closes. Each of these costs is modest in isolation; together they eliminate the vast majority of price-level signals.

Problem. Standard benchmarks evaluate stablecoin dislocation models on price-signal classification metrics (AUROC, F1). These metrics are blind to execution feasibility: a model that classifies price dislocations with 80% accuracy may still lose money on every trade it takes, because the dislocations it identifies are not executable. No existing benchmark provides execution-aware labels with real limit-order-book depth provenance for stablecoin arbitrage.

Contributions. We make four contributions:

- (1) Execution-aware benchmark. Real L2 order-book depth (Binance, Coinbase, Kraken), VWAP execution simulation, taker fees, and settlement penalties produce per-minute labels that distinguish price dislocations from executable dislocations.
- (2) Price-to-execution gap. Empirical measurement of the ratio between price-level and execution-level opportunity frequency during the SVB stress event: 12× at 1-minute horizon.
- (3) Oracle-gap evaluation. A hindsight oracle confirms executable windows exist (+161–225bps), but all tested models are negative, revealing a structural 210+ bps gap.
- (4) Open benchmark. Dataset, labels, model stack, and evaluation harness are open-sourced for future model comparison.

2 RELATED WORK

Stablecoin de-pegs. Cruz et al. [1] document USDC price fragmentation during the SVB collapse but do not evaluate ex-ante predictability. Kwon et al. [2] analyse Terra/LUNA mechanics; we use May 2022 as our validation event. Griffin and Shams [3] study Tether flows; our contribution is microstructure execution, not flow manipulation.

Limits to arbitrage. Hautsch et al. [4] model blockchain-specific limits to arbitrage theoretically (settlement finality, gas costs, latency). Shleifer and Vishny [5] and Gromb and Vayanos [6] establish the general theory. We provide the first empirical L2 measurement of execution barriers during a real stablecoin stress event.

Market microstructure. Cont and Kukanov [7] develop optimal order-routing across fragmented markets, which motivates our VWAP depth-walk. Glosten and Milgrom [8] and Makarov and Schoar [9] provide the theoretical and empirical basis for adverse selection during stress.

Financial ML benchmarks. Saly-Kaufmann et al. [10] propose ML benchmarks for crypto microstructure. Our work is

Table 1: Dataset coverage by split.

Split	Event	Minutes	Pos. rate
Train	Calm control (3 windows)	20,125	0.00%
Validation	Terra/LUNA, May 2022	11,523	2.45%
Test	USDC/SVB, Mar 2023	15,839	2.88%
Total		47,487	

Positive rate = executable arb at \$10K, 5-min horizon.

complementary: stablecoin-specific, event-based, and execution-labelled. Adams et al. [11] characterise AMM liquidity during stress, informing our settlement-penalty modelling.

3 BENCHMARK DESIGN

3.1 Data Sources and Instruments

The gold-layer dataset contains 47,487 1-minute bars across three event windows (Table 1). Price data comes from the Binance Vision public archive (klines, aggTrades) and Coinbase REST candles. Order-book depth for executable labels comes from Binance USDM futures bookDepth snapshots and incremental updates. We track four instruments BTC-USDC, BTC-USDT, ETH-USDC, ETH-USDT across Binance, Coinbase, and Kraken.

Depth provenance guarantee. Each row carries a `depth_source` tag: `real_l2_snapshot`, `real_l2_incremental`, or `synthetic_kline`. Net-profit labels use only `real_l2_*` rows; synthetic kline depth is excluded from paper-grade computations and isolated to a separate flag column (`is_paper_grade_depth`).

3.2 Event-Based Split Design

Splits are defined by event identity, not by time alone (Figure ??). The train split contains three calm control windows (no stress event) that bracket the test period; the validation split uses the Terra/LUNA collapse (May 2022), a qualitatively similar but independent stress event; the test split uses the USDC/SVB de-peg (March 10–15, 2023).

This design has two guarantees: (i) no model trained on the train split has seen stress dynamics; (ii) threshold calibration on validation uses an independent stress event, not a held-out portion of the same episode. No lookahead: labels are constructed by `join_asof` at $t + h$ and formally verified.

3.3 Feature Sets

We define four nested feature sets from narrowest to broadest:

<code>price_only</code>	5 cross-quote basis columns
<code>price_plus_book</code>	+7 microstructure cols (spread, depth, imbalance)
<code>price_book_frag</code>	+3 cross-venue fragmentation cols
<code>price_book_settle</code>	+7 on-chain settlement proxies

The nesting structure isolates the marginal value of each information layer.

4 EXECUTION-AWARE LABELS

4.1 Cross-Quote Basis

The cross-quote USDC basis at time t is:

$$b_{\text{USDC}}(t) = 10,000 \times \frac{P_{\text{BTC-via-USDC}}(t) - P_{\text{BTC-direct}}(t)}{P_{\text{BTC-direct}}(t)} \quad [\text{bps}]$$

where $P_{\text{BTC-via-USDC}}$ is the effective BTC price obtained by converting $\text{USD} \rightarrow \text{USDC} \rightarrow \text{BTC}$ relative to the direct BTC-USD market. We also compute the max-absolute basis $b_{\text{max}}(t) = \max(|b_{\text{USDC}}|, |b_{\text{USDT}}|)$ as a broader stress detector. Basis labels for horizon $h \in \{1, 5, 15, 60\}\text{m}$ and threshold τ :

$$y_{\text{basis}}(t) = \mathbf{1}[|b_{\text{USDC}}(t+h)| > \tau]$$

Primary classification task: `basis_usdc_1m_gt10bps`.

4.2 VWAP Execution Walk

For each minute window and notional $q \in \{\$10\text{K}, \$50\text{K}, \$100\text{K}, \$500\text{K}\}$, we simulate a round-trip execution:

- (1) Walk the real-L2 bid and ask ladders across venues to fill q notional, recording the VWAP price at each level consumed.
- (2) Compute the gross spread: $\text{gross}(q) = \text{VWAP}_{\text{sell}} - \text{VWAP}_{\text{buy}}$.
- (3) Deduct taker fees on both legs and a settlement delay penalty.

4.3 Net Executable Arbitrage Labels

$$\text{net}(q, t) = \text{gross}(q, t) - f_{\text{buy}} - f_{\text{sell}} - \delta_{\text{settle}} \quad (1)$$

$$y_{\text{arb}}(q, h, t) = \mathbf{1}\left[\max_{s \in [t, t+h]} \text{net}(q, s) > 0\right] \quad (2)$$

where $f_{\text{buy}}, f_{\text{sell}}$ are venue taker fees and δ_{settle} is a settlement penalty proxy. Primary economic task: `executable_arb_q10000_5m`.

The forward-max in Eq. (2) asks: does a profitable window exist within the next h minutes? matching the practical constraint that a trader can execute at any point within the horizon.

4.4 Oracle Upper Bound

The oracle trades every minute where $\text{net}(q, t) > 0$ in hindsight. It uses information unavailable at prediction time and sets the theoretical performance ceiling. The oracle is not deployable; it exists to confirm that profitable windows exist and to calibrate the oracle gap.

5 MODELS AND EVALUATION

5.1 Model Stack

Rule baselines. `NoTrade` (never trades; the economic anchor), `PriceBasis10bps` (trade when $|b_{\text{USDC}}| > 10$), `PriceBasis25bps`, and `GrossArbThreshold` (trade when gross spread > 0).

Statistical baselines. `LastValue` (AR0), `RollingMean`, AR1.

ML classifiers. Logistic (L2), Lasso (L1), Random Forest, XGBoost, LightGBM each trained on each of the four nested feature sets.

Table 2: Price-to-execution gap (test split, \$10K notional, 1-minute same-window horizon).

Thresh.	Price (max)	Price (USDC)	Exec. (\$10K)	Ratio
0 bps	97.15%	83.69%	3.34%	29×
5 bps	63.22%	25.77%	3.03%	21×
10 bps	35.09%	12.65%	2.88%	12×
25 bps	9.63%	6.53%	2.62%	3.7×

Expected-net-profit regressor. ExpectedNetProfitRegressor directly predicts $\text{net}(q, t)$ rather than the binary label, and trades when $\hat{y} > \theta^*$ where θ^* is calibrated on validation. This targets the economic objective directly without requiring a probability estimate.

5.2 Threshold Calibration

For probabilistic classifiers, the decision threshold is calibrated on the validation split:

$$\theta^* = \arg \max_{\theta \in [0.05, 0.95]} \sum_{i: \hat{p}_i > \theta} \text{net}(q, t_i) \quad \text{s.t.} \quad |\{i : \hat{p}_i > \theta\}| \geq 25$$

The ≥ 25 trade minimum prevents degenerate abstention strategies from winning on small sample artefacts.

5.3 Evaluation Protocol

Primary (economic): `net_bps_captured` (mean net bps per trade), `hit_rate_above_cost` (fraction of trades with positive net outcome), `final_pnl_usd` (cumulative P&L), `oracle_capture_pct` (= net bps / oracle net bps).

Secondary (statistical): AUROC, AUPRC, Brier score, n_{trades} .

The NoTrade baseline (= 0 net bps) is the economic floor: a model that loses money is worse than abstaining entirely, regardless of its AUROC.

6 RESULTS

6.1 Price-to-Execution Gap

Table 2 and Figure 1 show the central finding. At a 10 bps threshold, 35.09% of test-split minutes show a primary/max cross-quote basis exceeding the threshold on price alone; only 2.88% remain profitable after execution costs at \$10K notional a 12× price-to-execution ratio. The USDC-specific basis fires in 12.65% of minutes, reflecting that USDT routes also fragment during the SVB event. The gap persists at every threshold level.

6.2 Signal Waterfall

Figure ?? traces where value is lost at each stage of the pipeline:

Total test minutes	15,839	(100%)
Price disloc. > 10 bps	5,560	(35.1%, max basis)
of which USDC-specific	2,004	(12.65%)
Executable at \$10K / 1m	456	(2.88%)
Oracle profitable trades	351	(2.22%)

The oracle trades 351 minutes profitably; every model trades far more minutes but loses money on nearly all of them.




Figure 1: Price-to-execution gap by threshold. Price-signal frequency (dark) dwarfs executable frequency (light) at every level. The gap is largest at low thresholds, where price signals fire constantly but depth is thin.

Table 3: Oracle gap (test split). All non-oracle models are economically negative.

Task	Oracle	Best model	Gap
basis_usdc_1m_gt10bps	+161.7 bps	−49.1 bps	211 bps
executable_arb_q10000_5m	+224.6 bps	−42.9 bps	268 bps
executable_arb_q50000_5m	+146.8 bps	−112.7 bps	260 bps

Best model: `logistic@price_plus_book` (basis task);

`price_threshold_25bps@price_only` (arb tasks).




Figure 2: Oracle gap: grouped bars showing oracle net bps (gold) vs. best ML model (navy) for each task. The gap exceeds 200 bps across all tasks.

6.3 Oracle Gap

Table 3 gives the oracle gap for each primary task. The oracle earns +161.7bps on the `basis_usdc_1m_gt10bps` task. The best ML model (`logistic`, `price_plus_book` features) achieves −49.1 bps 211 bps below the oracle. On the executable arbitrage task, the oracle earns +224.6bps; the best rule (`price_threshold_25bps`) achieves −42.9 bps. Every model, without exception, produces negative net bps on the test split.

Table 4: Model ablation summary (task: **basis_usdc_1m_gt10bps**, test split). [†]Theoretical ceiling.

Model	Features	Net bps	Trades	AUROC
Oracle [†]		+161.7	316	
logistic	price_plus_book	-49.1	1	0.133
lgbm	price_only	-329.1	4	0.626
gross_arb	price_only	-136.0	611	0.530
price_thr 10	price_only	-269.5	2002	0.829
price_thr 25	price_only	-346.6	1033	0.743
no_trade		0.0	0	
ExpNetProfit	price_only	-61.4	1610	

6.4 Model Ablation

Table 4 summarises the model families. Price-threshold rules overtrade severely (611 trades, -136 bps): the rule fires on price dislocations with no executable depth. ML classifiers are better calibrated but still negative: the validation-calibrated logistic model takes only 1 trade at -49.1 bps. The Light-GBM model, despite higher AUROC (0.63), produces worse economic results (-329 bps) due to overtrading. Feature richness does not help: the best logistic result uses **price_plus_book**; adding fragmentation or settlement features does not improve economic outcomes.

6.5 Expected Net-Profit Regressor

The **ExpectedNetProfitRegressor** (LightGBM base, target: $\text{net}(q, t)$) achieves -61.4 bps on the test split better calibrated than rule baselines but still economically negative. This is the strongest evidence that the oracle gap is structural: even when the training objective directly matches the economic evaluation criterion, the model cannot identify profitable windows out of sample. The direct profit-prediction framing does not close the gap.

6.6 Robustness

Figure ?? and Table 5 show that the price-to-execution gap persists across all tested cost assumptions.

Fee regimes. At \$10K/5m/10 bps, base fees give 5.64% executable; high fees reduce this to 5.46%, raising the ratio from 2.24× to 2.32×. Institutional fees (lower) give 6.11%.

Settlement penalties. Adding 10 bps settlement penalty reduces executable windows from 5.64% to 4.88% (ratio: 2.24× → 2.59×). The combined worst case (high fee + 10 bps settlement) gives 4.81%.

Notional size. Smaller notionals fill more easily; at \$50K the executable rate drops to 1.62% (same window, 10 bps threshold). At \$500K, virtually no windows are executable (0.03%).

In all cases, the qualitative conclusion is identical: the gap is large, persistent, and model-invariant.

Table 5: Price-to-execution ratio across cost regimes (\$10K notional, 10 bps threshold, 5-min horizon, test split).

Settlement	Fee regime	Exec. %	Ratio
0 bps	low (institutional)	6.11%	2.07×
0 bps	base	5.64%	2.24×
0 bps	high	5.46%	2.32×
5 bps	base	5.16%	2.45×
10 bps	base	4.88%	2.59×
10 bps	high (worst case)	4.81%	2.63×

Table 6: False-positive diagnosis for **PriceBasis10bps** rule (test split).

Outcome	Count	%	Mean net bps
True positive (TP)	1,421	9.0%	+
False positive (FP)	581	3.7%	-38.7
False negative (FN)	583	3.7%	
True negative (TN)	10,672	67.4%	
Total test minutes	15,839	100%	

price > 10bps but not executable. Net bps shown for traded windows only.

6.7 False-Positive Diagnosis

Applying the **PriceBasis10bps** rule to the test split produces 2,002 trades. Table 6 decomposes the outcomes. False positives (581 windows: basis > 10 bps but no executable depth) account for 29% of trades and have a mean net profit of -38.7 bps, explaining the severe economic drag. The model fires on price dislocations where the order book is too thin to fill the arbitrage leg exactly the execution-barrier problem the benchmark is designed to measure.

7 LIMITATIONS AND FUTURE WORK

Single primary stress event. The test split covers one event (SVB March 2023). The Terra/LUNA validation event provides some cross-event evidence, but future work should extend to the USDC recovery period (March 15 – April 1, 2023) and the USDT/Curve stress event (June 2023).

Settlement proxy. CEX-to-CEX settlement latency is proxied via a fixed penalty parameter; true on-chain settlement delay is not modelled. Prokopenko et al. [12] provide a more detailed treatment.

Partial fills and sub-minute latency. The VWAP walk assumes all-or-nothing fills at a 1-minute snapshot. In practice, partial fills and sub-minute price impact change the effective execution cost.

Uncertainty-aware abstention. The uncertainty module (bootstrap ensemble, quantile regression) is implemented but its ablation is deferred to future work due to computational cost. Preliminary analysis suggests stricter abstention thresholds reduce false positives but do not achieve positive net bps.

Threshold ablation. The validation calibration uses the total P&L rule with a 25-trade minimum. A full ablation over fixed-threshold rules (0.5, 0.7), validation F1, and mean-bps objectives is planned to confirm that the negative result is not threshold-rule- dependent.

Feature scope. Order-flow imbalance at sub-minute resolution, funding rates, and on-chain stablecoin minting/redemption flows are absent from the current feature set.

8 CONCLUSION

Stablecoin StressBench introduces a transaction-cost-aware benchmark that separates price-level and execution-level stablecoin dislocations. The central finding is a 12× price-to-execution gap during the USDC/SVB crisis: 35.1% of minutes show a cross-quote basis exceeding 10 bps, yet only 2.88% survive a realistic VWAP execution filter at \$10K notional.

A hindsight oracle confirms the profitable windows are real (+161.7 bps average per trade), but every tested ML model, rule-based strategy, and direct net-profit regressor produces negative net bps out of sample an oracle gap exceeding 210 bps. Adding microstructure or fragmentation features does not close the gap; directly optimising the economic objective (ExpectedNetProfitRegressor) also fails (−61.4 bps).

The benchmark establishes stablecoin arbitrage prediction as an execution-barrier problem rather than a signal-detection problem. The challenge is not identifying when the price chart shows an opportunity, but predicting which chart opportunities will survive the order book.

ACKNOWLEDGEMENTS

The author thanks the Columbia University MAFN programme for supporting this research.

REFERENCES

- [1] Cruz, A., et al. (2024). USDC fragmentation during the SVB collapse: Cross-venue price gaps and liquidity withdrawal. *Journal of Financial Economics* (forthcoming).
- [2] Kwon, S., Minegishi, K., and Nishi, R. (2023). Terra/LUNA depeg mechanics: Algorithmic stablecoin collapse and contagion dynamics. Working Paper.
- [3] Griffin, J. M. and Shams, A. (2020). Is Bitcoin really untethered? *Journal of Finance*, 75(4):1913–1964.
- [4] Hautsch, N., Scheuch, C., and Voigt, S. (2018). Limits to arbitrage in markets with stochastic settlement latency. Working Paper, Vienna Graduate School of Finance.
- [5] Shleifer, A. and Vishny, R. W. (1997). The limits of arbitrage. *Journal of Finance*, 52(1):35–55.
- [6] Gromb, D. and Vayanos, D. (2010). Limits of arbitrage: The state of the theory. *Annual Review of Financial Economics*, 2:251–275.
- [7] Cont, R. and Kukanov, A. (2013). Optimal order placement in limit order markets. *Quantitative Finance*, 17(1):21–39.
- [8] Glosten, L. R. and Milgrom, P. R. (1985). Bid, ask and transaction prices in a specialist market with heterogeneously informed traders. *Journal of Financial Economics*, 14(1):71–100.
- [9] Makarov, I. and Schoar, A. (2020). Trading and arbitrage in cryptocurrency markets. *Journal of Financial Economics*, 135(2):293–319.
- [10] Salý-Kaufmann, et al. (2026). ML benchmarks for cryptocurrency market microstructure. *Proceedings of ICAIF 2026*.
- [11] Adams, H., et al. (2021). Uniswap v3 core. Technical Report, Uniswap Labs.
- [12] Prokopenko, A., et al. (2023). On-chain settlement latency under stress: Empirical evidence from Ethereum mainnet. Working Paper.
- [13] He, Z. and Krishnamurthy, A. (2013). Intermediary asset pricing. *American Economic Review*, 103(2):732–770.
- [14] Gorton, G. and Pennacchi, G. (1990). Financial intermediaries and liquidity creation. *Journal of Finance*, 45(1):49–71.
- [15] Pennington, J., et al. (2022). Failure modes of algorithmic stablecoins. Working Paper.